



Qian Jiang



(+86) 182 1144 5116 | 202316187@stu.neu.edu.cn | Huizhou and Shenzhen, Guangdong, China

Education

Peking University, Shenzhen Graduate School Project 985 · Project 211 · Double First-Class

2027.09 - 2030.06

M.S. in Computer Science

- Admitted

Northeastern University Project 985 · Project 211 · Double First-Class

2023.09 - 2027.06

B.Eng. in Computer Science and Technology

- GPA: 4.19/5.00 (5/196, top 2.56%)
- CET-4: 579, CET-6: 528
- Selected coursework: Calculus I/II (99), Probability Theory (100), Optimization (97), Linear Algebra (92)

Publications

- Four publications and 80 Google Scholar citations: one CCF-A paper as co-first author, one CCF-B paper as first author, and two CCF-A papers as a contributing author (* denotes equal contribution).

1. **(Co-first author, CCF-A, ACL 2026 Main)** Zirui Song*, Qian Jiang*, et al. Audio Jailbreak: An Open Comprehensive Benchmark for Jailbreaking Large Audio-Language Models.
2. **(First author, CCF-B, ICME 2026 Oral)** Qian Jiang, Zhecheng Shi, et al. OmniFood-Bench: Evaluating VLMs for Nutrient Reasoning and Personalized Health Advice.
3. **(Contributing author, CCF-A, AACL 2026)** Zirui Song*, Guanghan Ouyang*, ... Qian Jiang ..., et al. ManipLVM-R1: Reinforcement Learning for Reasoning in Embodied Manipulation with Large Vision-Language Models.
4. **(Contributing author, CCF-A, ACL 2026 Findings)** Fan Ji*, Juncheng Yang*, ... Qian Jiang ..., et al. FineState-Bench: Benchmarking State-Conditioned Grounding for Fine-grained GUI State Setting.

Research Experience

Mohamed bin Zayed University of Artificial Intelligence, advised by Prof. Xiuying Chen

- Conduct research on trustworthy AI, including multimodal LLM jailbreak attacks and reliable AI-generated text detection.

Language and Intelligent Systems Lab, Northeastern University, advised by Prof. Miao Fang

- Conduct research on VLMs and agents; serve as the lead student coordinator for laboratory recruiting and operations.

Selected Research Projects

Audio Jailbreak: An Open Comprehensive Benchmark for Jailbreaking Large Audio-Language Models — Co-first Author, CCF-A, ACL 2026 Main

Built AJailBench, an open jailbreak benchmark for large audio-language models, filling a gap in audio-model safety evaluation. Curated 1,495 high-quality adversarial samples from tens of thousands of candidates across ten violation categories and built a unified pipeline for seven leading audio-language models, revealing that no model is robust across all dimensions. Proposed APT, a toolbox of ten temporal, frequency-domain, and hybrid perturbations with strict semantic-consistency constraints; Bayesian optimization used a cosine-similarity objective to identify jailbreak-prone parameters and improved attack success by about 10%. The code was adopted by the official MBZUAI team, downloaded 784 times on Hugging Face in one month, invited for an AI Bench Talk, and reported by media including MSN. Accepted to ACL 2026 Main with me as a co-first author.

OmniFood-Bench: Evaluating VLMs for Nutrient Reasoning and Personalized Health Advice — First Author, CCF-B, ICME 2026 Oral

Addressed the limitations of food benchmarks that focus on coarse-grained classification and omit safety reasoning by building a three-stage VLM evaluation framework: basic perception, quantitative reasoning, and safety-critical advice. Constructed 1,208 high-quality MM-Food-100K samples across home, restaurant, packaged-food, and fresh-food settings. Evaluation of six SOTA VLMs, including GPT-5.1 and Gemini-3-flash, revealed a semantic-physical gap: near-human dish recognition but catastrophic mass-estimation failures and frequent safety hallucinations in high-risk diabetes and kidney-disease scenarios. The best safety-advice accuracy was only 46%, providing a rigorous basis for trustworthy medical-grade dietary models. Accepted as an ICME 2026 Oral paper with me as first author.

AI Text Detectors as a House of Cards: From Vulnerability Induced by Decoding Strategies to Robustness Through Restoration

Found that existing logit-based AI-text detectors are extremely vulnerable to decoding strategies: small changes in temperature or repetition penalty can reduce detection AUC to about 50%. The root cause is that decoding changes token-probability distributions, making human and machine text features overlap. Derived a unified theoretical explanation for major detection methods and proposed Restore, an adaptive distribution calibrator that learns mappings from features including Distinct-2, trigram repetition ratio, and character entropy. Restore maps shifted distributions back into a detectable space and substantially recovers detector robustness across models and datasets.

Honors & Activities

- 2026 Mathematical Contest in Modeling Finalist (top 0.1%-2%)
- Third Prize, National Final of the Global AI Algorithm Elite Competition; proposed the SLTP travel-planning framework with dual-expert LoRA and teacher-student distillation for accurate symbolic conversion by lightweight models, plus heuristic tree search balancing constraint compliance and solution diversity.
- Technical blog: <https://www.zhihu.com/people/bu-neng-yao-gong>
- Student lead, undergraduate academic community: <https://neuqnlp.github.io/team>